Response revision working note

Prior manuscript: TNSM-2026-10776

14 November 2026

This working note supersedes the old letter's positive DeadlineGuard and load-generalization interpretation. It is not a declaration that all review comments or independent acceptance checks are closed. The prior full letter is historical source material; its unchanged excerpts must not be submitted with this note as if current.

Reviewer 2, Comment 1: value beyond a static detector

Reviewer concern

The reviewer asks whether dynamic scheduling beats the simplest static alternative, and requests changing-traffic conditions rather than unnecessary complexity.

Author response

We agree and now distinguish static operating-point selection from adaptation. A corrected follow-up retains the six-method order, frozen classifiers and threshold 0.5, with ten blocks of the fixed load–prevalence factorial. Both Q and DQN select LightLR throughout, while DeadlineGuard selects MedRF throughout; their detection summaries match the corresponding static baselines. All methods have zero recorded deadline-miss windows. We therefore withdraw the earlier use of DeadlineGuard switching superiority as the new study's positive result. Q and DQN show no resolved energy benefit over directly deploying LightLR; intervals crossing zero do not prove equivalence. The fixed confidence cascade instead has a specific energy–detection trade-off.

Location

Current manuscript: corrected state/system, objective/training support, follow-up results, discussion and limitations; all three sensitivity sets are reported in the follow-up tables.

Revised interpretation

DQN minus Static-LightLR is $+0.599$ J, 95% paired CI $[-4.119, 5.317]$ J; Q minus Static-LightLR is $+0.246$ J $[-4.055, 4.547]$ J. Cascade minus Static-MedRF is -12.064 J $[-16.548, -7.580]$ J, at -0.0446 F1 percentage points. These are not claims of general RL uselessness or cost-free cascade improvement.

Reward, state and information concerns

Author response

The common view now uses current arrivals, FIFO carryover before appending arrivals, prior selected-detector threat and previous action. No temperature reward is fabricated. Training includes deadline and pending-work penalties, but the calibrated fixed surrogate never develops backlog, so no queue-adaptation competence or queue-ablation robustness is claimed. Fifteen historical constant-coordinate ablations do not validate changes in those coordinates. Immediate pairwise utility reversals do not establish long-horizon or global policy optimality. Cascade's current-row confidence is observed after LightLR entry inference, not at the RL pre-window decision. These computation and training-coverage differences are disclosed rather than labelled identical comparisons.

Measurement, provenance and independent-review scope

Author response

All sixty accepted runs and both invalid attempts remain archived. Blocks 1 and 5 span explicitly authorized recovery periods. Raw and idle-adjusted results are reported for all blocks, exclusion of block 1, and exclusion of both recovered blocks. The monitor cadence repair and rejected-snapshot logging preserve the 5% external CPU and 1.5-second gap limits; exact old rejected-gap duration is unavailable. Historical results are never pooled or overwritten. Author-side validation and JSON recomputation are separate from independent ZIP/raw-log acceptance, which remains incomplete due to the review environment. We do not mark that boundary as closed or request further physical measurements.